

YCSB vs BigBench

Berkay Barlas

Technical University of Munich
berkay.barlas@tum.de

Achraf Aroua

Technical University of Munich
achraf.aroua@tum.de

ABSTRACT

With the increase of data volume, velocity, and variety, we have seen the emergence of many systems that supports big data storage and data processing. Due to their different architectures, it has become necessary to evaluate and compare these systems.

In this paper, we provide an overview of Yahoo! Cloud Serving Benchmark(YCSB) and BigBench, two big data systems benchmarks. We describe their design goals, workloads, and their architecture.

KEYWORDS

Benchmark, YCSB, BigBench, Workloads, Design goals, Architecture

ACM Reference Format:

Berkay Barlas and Achraf Aroua. 2020. YCSB vs BigBench. In *Proceedings of TUM*. ACM, New York, NY, USA, 6 pages. <https://doi.org/isds>

1 INTRODUCTION

The increasing data *volume*, the requirement of quick data processing(*velocity*), and the wide variety of data types(*variety*) give various problems to traditional relational database management systems which are unable to solve. This led to the explosion of new systems with the goal of big data storing and processing. Despite their common goal, different decisions were made when implementing the systems to solve various problems faced in big data systems. For example, some systems are more optimized for writes while others have decided to optimize for reads. The absence of a standard database benchmark caused database management systems developers to make performance claims based on self-defined and biased benchmarks. These specifically built benchmark

tools are used to make misleading comparisons and claims. Therefore, benchmarks such as YCSB[7] and BigBench[1] emerged to tackle these issues and solve these problems by bringing standardized benchmark tools.

The paper is organized as follows. Section 2 provides an overview of the design goals of YCSB and BigBench. Section 3 discusses the workloads of each benchmark in detail, while section 4 examines their architecture. Furthermore, section 5 gives a comparison of the two benchmarks. Section 6 examines further work. Lastly, section 7 presents our conclusions.

2 DESIGN GOALS

The Yahoo! cloud serving Benchmark was first presented in 2010 to evaluate and compare the performances of different cloud serving systems. These cloud systems are usually NoSQL systems that do not support ACID criteria. YCSB is used to evaluate the performance of cloud systems by examining their scalability and processing performance.

BigBench, the first end-to-end benchmark tool for big data system, was first published in 2013. It is based on a fictional product retailer. BigBench focuses on parallel database management systems and map reduce engines. Also, BigBench V2, an new version of BigBench with improved workload queries and data generator, published in 2017.

2.1 YCSB Tiers

In this section, we will see the two different tiers of the YCSB benchmark that aims at evaluating the performance and the scalability of cloud serving systems

2.1.1 Tier1:Performance:

Since cloud serving systems aim for supporting very high request rates, we try in this tier to see the latency of the system when it is under load. It is expected that the latency of the requests to increase when the number of throughput (the number of operations per second) increases. This can be mostly explained to the less space for disk, CPU, and so on. Thus, there is a trade-off between latency and throughput.

This performance tier helps in examining this trade-off for different database systems by measuring the latency as we increase the throughput until the system is saturated and throughput stops increasing

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

TUM, 2020,

© 2020 Association for Computing Machinery.
ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00
<https://doi.org/isds>

2.1.2 Tier2:Scaling:

YCSB uses two metrics when evaluating the scaling aspect of different cloud serving systems:

- **Scale-up**: Here, we examine the change and the impact on performance when machines are added to the system. This can be easily checked when doing this following experiment:

Part-1

- Step 1: Load data to a certain number of servers
- Step 2: Run the workload

We try to save the results of this workload to use it later to compare it with the latter results.

Part 2

- Step 1: Delete data from servers
- Step 2: Add more servers
- Step 3: Load more large data
- Step 4: Run the workload

In the end, we compare the results of this workload with the first one. If a system has good scaling properties, we will get constant performance. More importantly, to make this experiment fair, the number of server and amount of data should increase proportionally

- **Elastic speedup**: This metric is quite similar to the first scale-up metric. However, the key difference is to check the performance of databases when machines are added but while the system is running. This can also be checked with the following experiment :

- Step 1: Load data to a certain number of servers
- Step 2: Run the workload
- Step 3: Increase the number of servers

after we increased the number of servers, we can conclude if the system offers the elasticity properties. If there's a performance improvement in a short time, then this system is considered as elastic. Usually, Systems can not show this improvement instantly because it needs time to reconfigure itself due to the newly added servers.

2.2 BigBench Design Goals

As the number of big data applications increased different requirements arose which led to the development of different database management systems. Therefore, a standard comparison and benchmark system are required to evaluate different big data systems.

BigBench is designed as end to end benchmarking tool that covers all the events during the initialization and operation of big data system.

BigBench is created in cooperation with the University of Toronto, Oracle, Teradata, and Infosizing. It is adopted from

another benchmarking tool TPC -DS which developed by Transaction Processing Performance Council (TPC)[5].

BigBench is based on a fictional product retailer system. Therefore, BigBench initializes and operates the database system as it is a product retailer who sells products to customers via physical and online stores

3 WORKLOADS

3.1 YCSB Workloads

To conduct the tiers presented in the previous section, YCSB comes with several built-in workloads. A workload represents a particular mix of read/write operations, data sizes, request distributions, and so on. YCSB has a package called **Core Package** which is a collection of related workloads. Each of these workloads tries to examine and evaluate systems in a particular performance space. These workloads do not cover the entire performance space but the most relevant ones to a cloud serving system. Since cloud systems have some trade-offs like synchronous vs asynchronous reads, these workloads aim at evaluating them. However, the user can also define new workloads thanks to the extensibility of YCSB. This process is defined in section 4.

Despite their different properties, these workloads are just a variation of the same basic application. In this application, there is a table of records, each with F fields. Each record is identified just by a primary key like "user145143". The values of each field are random ASCII characters of length L. Then for each application, it exists L * F byte records at the total. Against these data, each operation is randomly chosen to be one of: Insert, Update, Read, Scan. Since YCSB deals only with cloud serving systems, the YCSB package contains different combinations of workloads that aims to model different cloud systems application. These workloads are shown below in **Table 1**

Table 1: Different workloads in the Core Package

Workload	Operations	Record selection
A- Update heavy	Read:50%, Update:50%	Zipfian
B-Read heavy	Read:95%, Update:5%	Zipfian
C-Read only	Read:100%	Zipfian
D-Read latest	Read: 95%, Insert:5%	Latest
E-Short range	Scan:95%, Insert:5%	Zipfian/Uniform

3.1.1 YCSB Distributions.

Based on **Table 1**, the behavior of selecting a record is different from one workload to another. When running a workload,

various random choices need to be done: which record to run, how many records to scan, which operation to perform, and many others. Thus, each of these workloads is associated with some random distribution. Fortunately, YCSB has several probability distributions :

- **Uniform:** All random action have the same constant probability. An example would be when choosing an operation, is that all operations have the same probability to be performed
- **Zipfian:** In this distribution, some items will be extremely popular(the head of distribution) while others will be usually unpopular(the tail of the distribution)
- **Latest:** This distribution is a special case of the zipfian where the most inserted items are the head of distribution
- **Multinomial:** Probabilities for each item are specified and assigned at the beginning. Thus each item has a constant specific probability to be chosen

Thanks to these different distributions, different applications behavior can now be mimicked by the workloads which can help us in benchmarking the systems.

3.2 BigBench Workloads

BigBench workload depends on 30 queries. And these 30 queries designed to cover business and technical functions. These workloads defines the queries that will be used to test data system.

3.2.1 Business Dimension.

For business functions these workloads are designed cover the different categories of big data analytics proposed by McKinsey[4]. BigBench covers 10 big data levers in 5 business category.

- (1) Marketing
 - Cross-selling
 - Customer micro-segmentation
 - Sentiment analysis
 - Enhancing multi-channel consumer experience
- (2) Merchandising
 - Assortment optimization
 - Pricing optimization
- (3) Operations
 - Performance transparency
 - Return analysis
- (4) Supply chain
 - Inventory management
- (5) New Business Models
 - Price Comparison
 - The individual entries are indicated with a black dot, a so-called bullet.
 - The text in the entries may be of any length.

3.2.2 Technical Dimension.

For technical functions, these workloads are designed to span three different dimensions based on data sources, query processing types and analytic techniques.

(1) Query Processing Type

This technical function focuses on processing type of queries; declarative and procedural in addition to mixture of both. This means some queries are declarative can be used with SQL or procedural can be used with Map-Reduce while others are can be used with mix of both with a language like Pig Latin.

(2) Data source dimension

BigBench targets three different data types; structured, semi-structured and un-structured. Queries on web click events checks semi-structured and queries about product reviews check un-structured data. There are also some queries that works on multiple data types.

(3) Analytic Techniques

In this functionality BigBench workloads aims different business analytic methods. The three main areas used by defining queries are simple reporting, data mining and statistical analysis.

To summarize the workload queries in BigBench are defined according to business cases in a way that covers technical dimensions.

4 ARCHITECTURE

4.1 YCSB Client Architecture

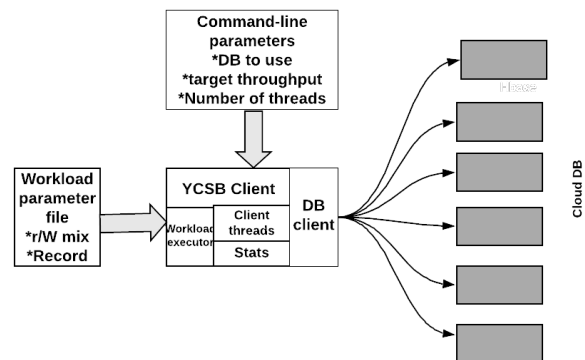


Figure 1: YCSB client Architecture

The YCSB client written in Java is a workload and data generator. It is responsible for generating data to be loaded to the database and the operations to be executed on them. The client has different modules and its architecture can be seen in Figure 1. The client has a command-line and a workload parameter file which are responsible for letting

the user pass the parameters to the client. These parameters represents **the workload** and **run-time properties** of the client. First, the workload properties are the read/write mix of the database, the distributions, and the number of records. These properties can be shared among different experiments. The run time properties are on the other hand are specific for a given experiment. This includes the database interface layer which is different for every system and the number of threads and so on.

Furthermore, multiple threads are running inside the client. Each of these threads is responsible for executing a series of operations by making calls to the database interface layer. These calls can happen in two phases: First in **the load phase** when loading the data to the database and in **the transaction phase** which corresponds to executing the workload. Moreover, these threads are also responsible for measuring the latency and the achieved throughput of the operations and report these results back to the statics module. Finally, the statics module aggregates the measurements and reports an average 95th and 99th% latencies. The results are returned as either a histogram or a time series of reported latencies.

4.2 YCSB Extensibility

Due to the various amount of different cloud serving systems, one of the key features of YCSB is that it is easily extensible. Therefore, developers can extend it to benchmark a cloud system different from the one originally supported by YCSB. As discussed in subsection 4.1 the database interface layer is responsible for translating requests from the client threads to the databases. Thus, you can extend the YCSB client in order to benchmark new database systems by writing a new class implementing the following methods: **read()**, **insert()**, **update()**, **delete()**, **scan()**. These represent the **CRUD** methods which allow us to work with a new database system. The user can also extend the client by defining new packages of new workloads in two ways. The most genuine way is to define a workload that uses **CoreWorkload** but with different parameters. By changing the parameters, the user can choose which operations to perform, which distribution to use and the number of records to include in the workload. The second approach which is the harder way is to write a new Workload executor Java class. With this approach, the user can perform complex operations and extend the workload to evaluate different points in the performance space.

4.3 BigBench's Architecture

BigBench data model d considers all important features of a big data system in terms of the three Vs, volume, variety, velocity described by Douglas Laney[3]:

4.3.1 BigBench's Data Model.

Bigbench is using an improved version TPC-DS data model. While the previous data model was only consisting of structured data, the new data model also includes a semi-structured data type for weblogs, and click stream events, and unstructured data type for product reviews. The improved data model is shown in Figure 2.

Their design assumes

- The semi-structured data is in a key-value format which is similar to the log format of Apache's web server.
- The un-structured data used by product reviews doesn't require users.

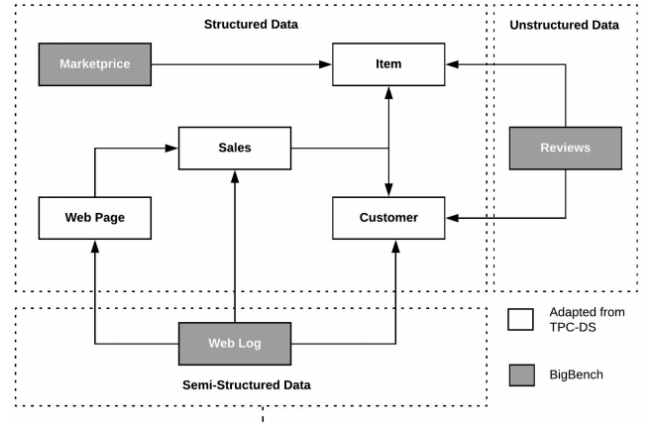


Figure 2: BigBench Data Model

4.3.2 BigBench's Data Generator.

To test a big data system end-2-end, BigBench generates required data for benchmarking for 3 different data types. The data generator used in BigBench is an extended version of The Parallel Data Generation Framework (PDGF)[2]. PDGF is a parallel data generator tool implemented in Java and fully platform-independent for producing large amounts of data. Previously, it was only generating structured datatype, while, improved version supports semi-structured and unstructured data types.

4.3.3 BigBench Review Generation.

Reviews are the unstructured type of data. To create a much more realistic data set they play a crucial role. BigBench uses an algorithm called TextGen for review reproduction. The review generator developed for BigBench is

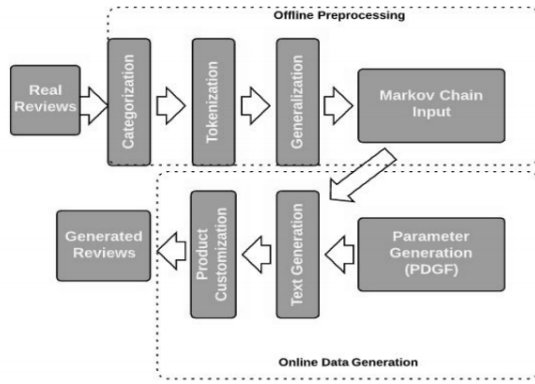


Figure 3: BigBench Review Generation

another independent tool that can be modified using XML. It produces text reviews based on samples. Their algorithm TextGen uses a Markov Chain method which creates a dictionary from the words in the given input. BigBench uses real product reviews from amazon.com for as initial sample data. Figure 3 shows the outline of the review production process.

The review generation consists of two phases

- (1) Reprocessing the data
- (2) Online data generation

The preprocessing part categorizes and prepares data to process with the Markov Chain algorithm. It creates a dictionary to be used in the review generation. The transition possibilities, Markov chains, of a word in the input reviews calculated. BigBench uses order-2 chains to generate product reviews. During the online data generation part, the generator uses products in the system to generate reviews according to produced parameters from PGDF.

4.3.4 BigBench Metrics and Evaluation.

BigBench doesn't have a final metrics design, on the other hand, it is suggested to use query execution times as a performance metric similar to TPC-DS. Instead of using specific queries using query processing type to categorize execution times will bring out the processing performance.

BigBench is tool that focuses on database management systems and Map-Reduce engines, thus, any of these systems or combination of them should be able used in benchmark. Figure 4 shows an example architecture of DBMS system that BigBench can evaluate.

BigBench begins evaluation with creation of tables. It uses DSDGEN from TPC-DS to create structured data, and improved version of PDGF to create new parts in data model; web logs and products reviews. The generated data loads into

tables with additional changes if required such as parsing web logs to create a schema. After loading the data, BigBench executes defined workload queries using an interface and execution times of each query are reported.

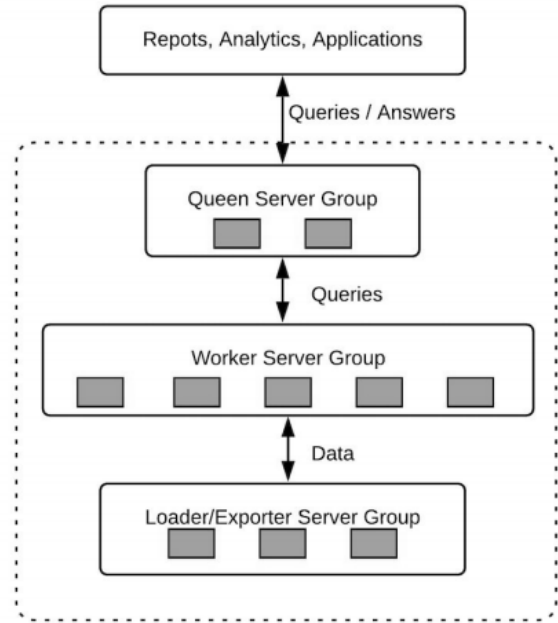


Figure 4: Example DMBS architecture for BigBench Evaluation

5 COMPARISON

YCSB and Bigbench are two benchmarking tools for big data systems. They are designed to create a standard comparison system.

Both of the benchmarking tools are open source and end-to-end benchmarks

However, YSCB and Bigbench have distinct features from each other which makes them useful for different situations. YCSB provides an apples to apples performance comparison of systems. It focuses on *serving* systems that have different tradeoffs like read vs write performance. Therefore, YCSB tries to examine them directly with workloads that are mixes of OLTP(Online Transactional Processing) operations. However, Bigbench is based on a fictional retailer. This means that

it uses certain workload queries to evaluate data management systems which makes it a better solution for complex comparisons needed for particular applications.

6 FURTHER WORK

6.1 YCSB++

YCSB++ first presented in 2011, is a set of extension to YCSB framework. This benchmark is used to evaluate the advanced features of NoSQL systems like Bulk loading, fast ingest techniques and so on. These extensions are *multi tester coordination*, *eventual consistency measurements*, *workloads multi phase workloads* and they allow for more complex tests. One important improvement would be parallel testing. In YCSB, at most one client can be run at once. However in YCSB++, we have a zookeeper-based synchronisation for multiple client allowing the different tests to be coordinated

6.2 BigBench V2

BigBench V2 is published in 2017 as an improved version to solve drawbacks of BigBench. The data generator used in first version improved to produce structured, semi-structured and unstructured data types according to a scale factor in new version. The data model changed to simple star schema from previous complex schema to keep up with recent big data systems. In addition to these improvements, BigBench v2 assumes web logs are in key value format instead of converting them to structured type.

7 CONCLUSION

In this paper, we explained the need for a tool that compares different cloud management systems and introduced two benchmarking tools for big data systems: YCSB and BigBench. After examining the metrics of each benchmark, this paper reviews their architecture and their workload generation technique. Additionally, by comparing the systems, we gave an overview of the two implemented improvements to the systems which are YCSB++ and BigBenchV2.

Over a decade YCSB became popular for benchmarking different cloud serving systems. Currently, YCSB is implemented and can be run with more than 18 different engines like *Cassandra*, *HBase*, *MongoDB*, etc.

BigBench also became one of the most known benchmarking tools due it's capabilities in realistic evaluation results. It's adopted as TPCx-BB by the Transaction Processing Performance Council.

REFERENCES

- [1] A. Ghazal, T. Rabl, M. Hu, F. Raab, M. Poess, A. Crolotte, and H.-A. Jacobsen, *BigBench: Towards An Industry Standard Benchmark for Big Data Analytics*. in SIGMOD, 2013, pp. 1197–1208.
- [2] T. Rabl, M. Frank, H. M. Sergieh, and H. Kosch. *A Data Generator for Cloud-Scale Benchmarking* In TPCTC, pages 41–56, 2010.
- [3] D. Laney. *3D Data Management: Controlling Data Volume, Velocity and Variety*. Technical report, Meta Group, 2001.
- [4] J. Manyika, M. Chui, B. Brown, J. Bughin, R. Dobbs, C. Roxburgh, and A. H. Byers. *Big data: The Next Frontier for Innovation, Competition, and Productivity*. Technical report, McKinsey Global Institute, 2011. http://www.mckinsey.com/insights/mgi/research/technology_and_innovation/big_data_the_next_frontier_for_innovation/
- [5] TPC Benchmark DS
- [6] Ghazal, Ahmad and Ivanov, Todor and Kostamaa, Pekka and Crolotte, Alain and Voong, Ryan and Al-Kateb, Mohammed and Ghazal, Waleed and Zicari, Roberto, *BigBench V2: The New and Improved BigBench* 10.1109/ICDE.2017.167, 2017, pp. 1225-1236.
- [7] Benchmarking Cloud Serving Systems with YCSB Yahoo! Research Santa Clara, CA, USA cooperb,silberst,etam,ramakris,sears@yahoo-inc.com 978-1-4503-0036-0/10/06 <https://www2.cs.duke.edu/courses/fall13/compsci590.4/>
- [8] YCSB++: Benchmarking and Performance Debugging Advanced Features in Scalable Table Stores. Swapnil Patil, Milo Polte, Kai Ren, Witawat Tantisiriroj, Lin Xiao, Julio Lopez, Garth Gibson, Adam Fuchs, Billie Rinaldi. Proc. of the 2nd ACM Symposium on Cloud Computing (SOCC '11), October 27–28, 2011, Cascais, Portugal. Supersedes Carnegie Mellon University Parallel Data Laboratory Technical Report CMU-PDL-11-111, August 2011. <https://www.pdl.cmu.edu/ycsb++/>
- [9] Benchmarking Big Data Systems: A Review IEEE TRANSACTIONS ON SERVICES COMPUTING, VOL. 11, NO. 3, MAY/JUNE 2018 University of Texas at Austin, Austin, TX 78712. Digital Object Identifier no. 10.1109/TSC.2017.2730882 <https://www.researchgate.net/publication/318664623>